





Counted in the representation actually trained (fp32), LoRA is more efficient. A ternary-reserve win is a PREDICTION contingent on the untested Gate 0 — genuinely open (the format assay measures from-scratch QAT, not a frozen-base reserve).