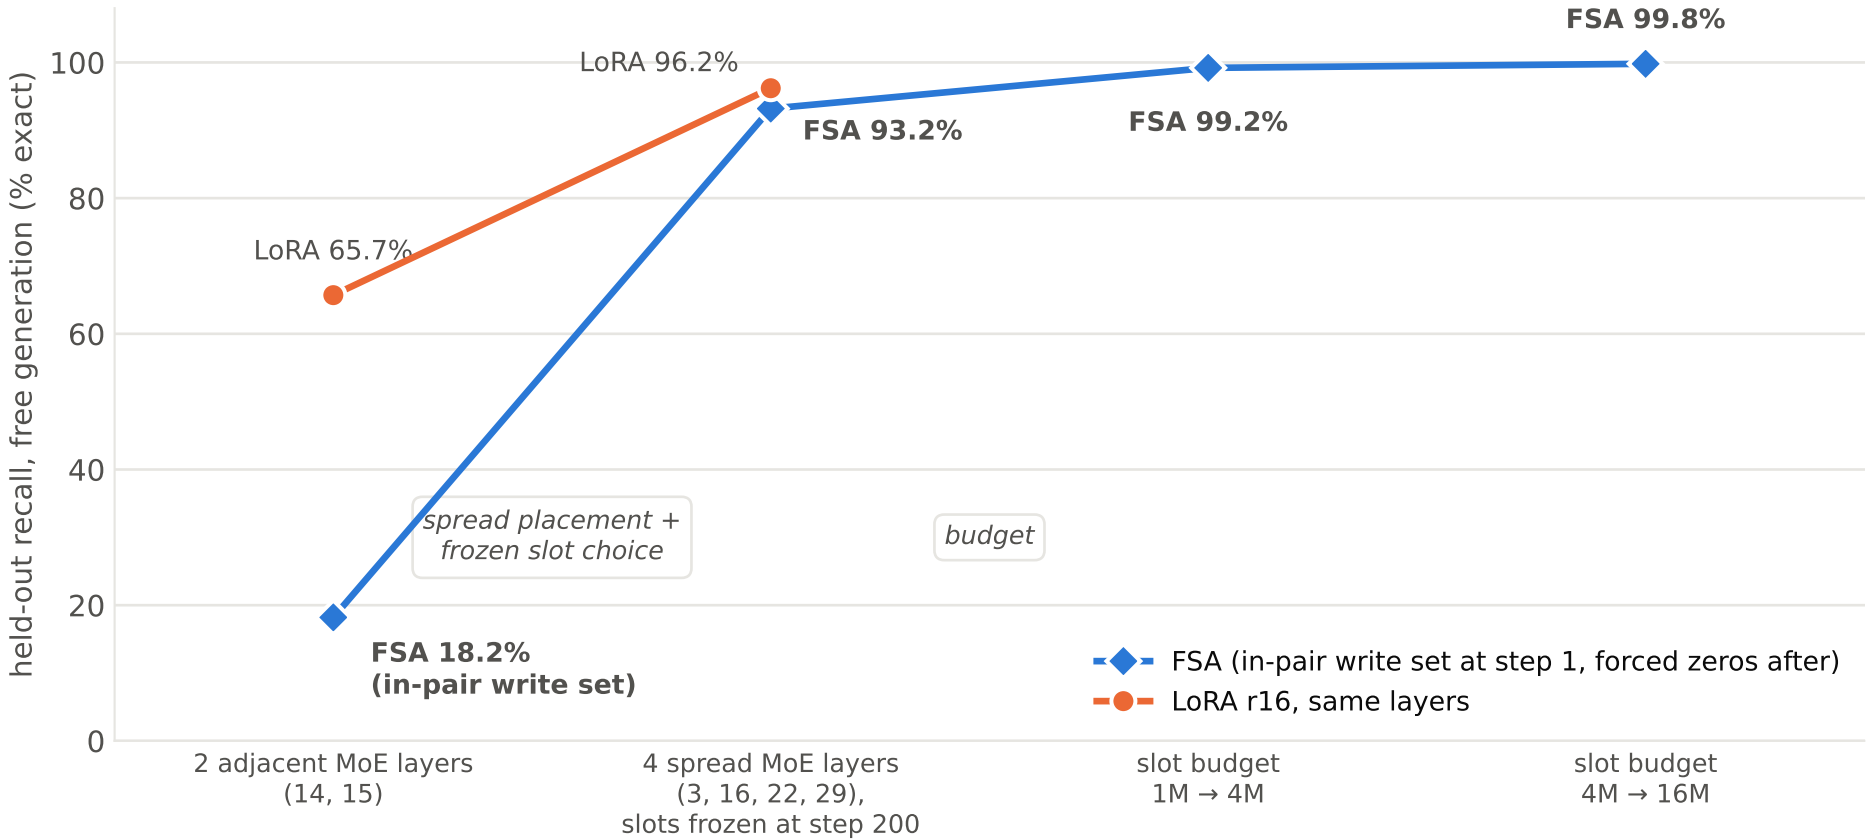

One seed per point (seed 1 for the 4-layer runs). Step 1 (2 layers): FSA used the earlier in-pair write set, LoRA r16 + penalty, both single-seed test runs; the steps do not isolate write set, placement, freezing and training length from one another. The 1M-budget recipe rerun on the current code scored 91.3% (vs 93.2%). Patches (packed, scale derived from the base; measured): 2.66 MB at 4M (code damage +0.011), 3.31 MB at 16M (2.13M slots written, code +0.007); both fail the +0.003 gate, as LoRA's +0.008 does. Every point is final-step test scoring. Not shown: the 1M recipe trained 800 steps instead of 400 scored 99.5% (one seed; the 400- and 800-step runs also diverge early for a reason not yet isolated).