Same facts, fewer values written: the best tested FSA recipe matches or beats LoRA r16 on three fact tasks

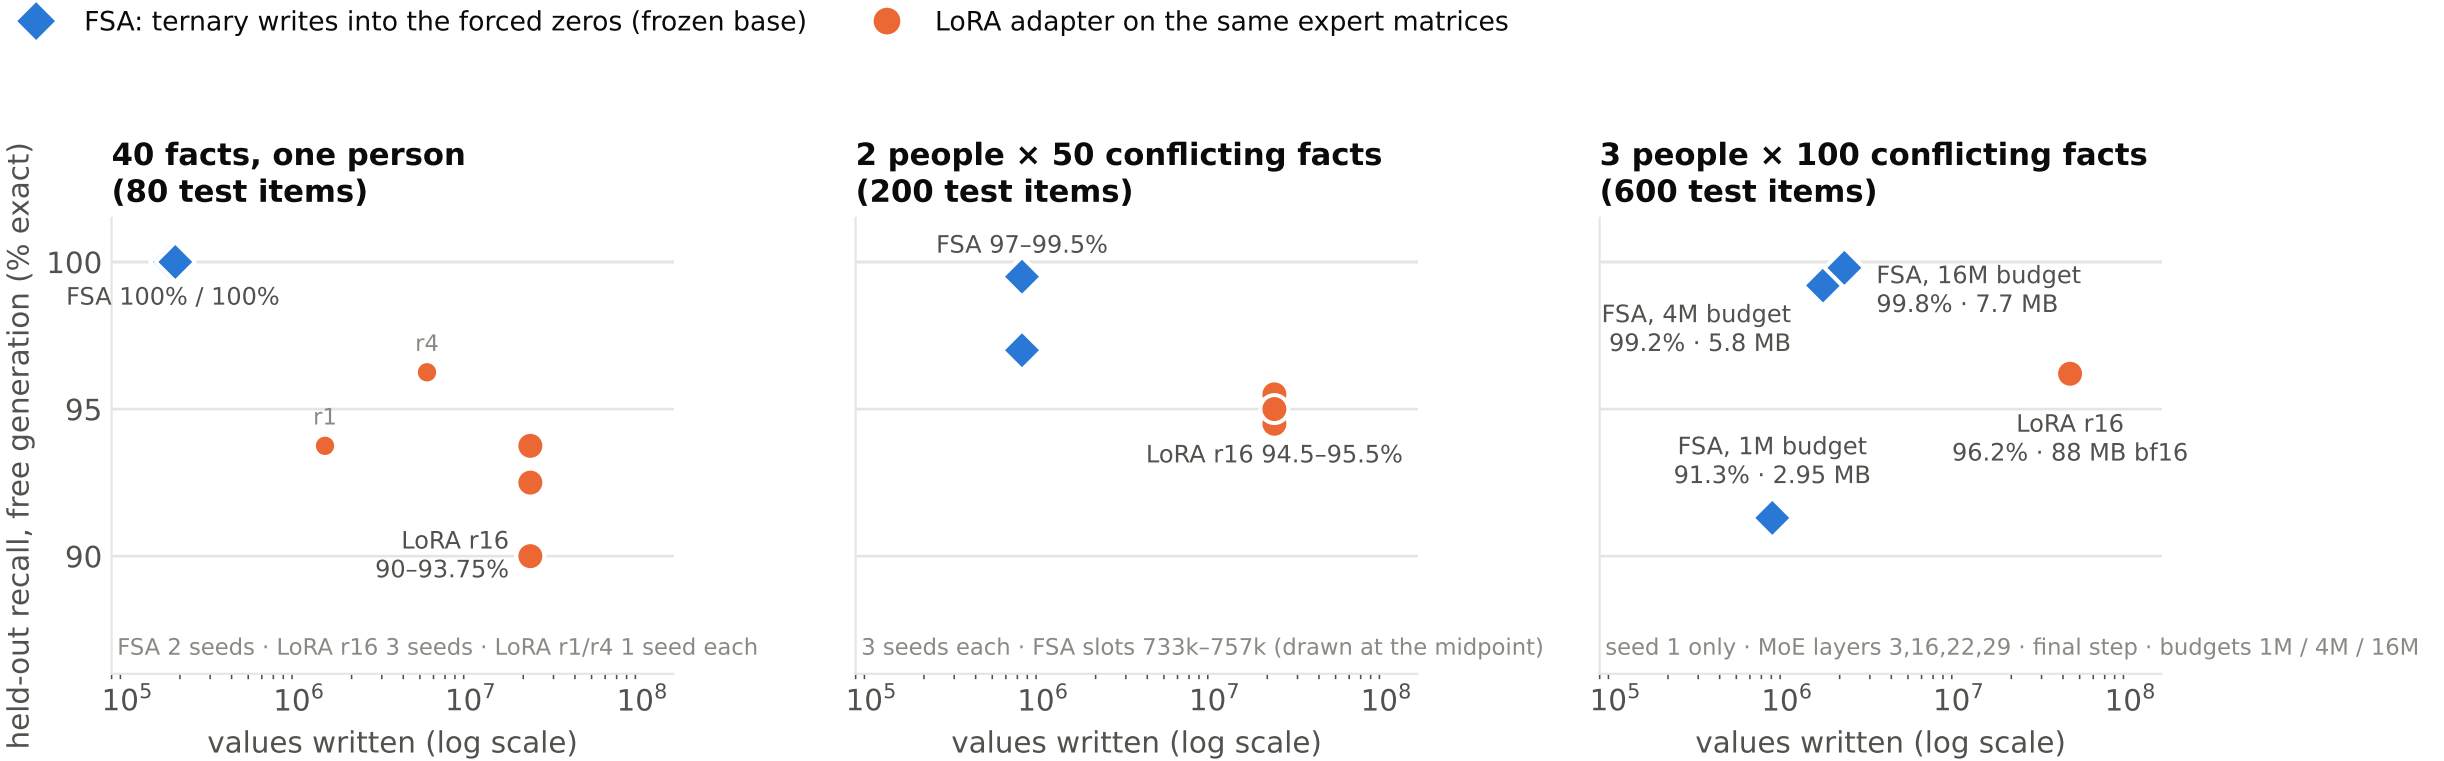

Test = held-out phrasings, greedy generation, whole answer must match. x: FSA = filled ternary slots, LoRA = adapter parameters (bf16 bytes = 2 × params; int4 factors would be 4× smaller). Lambda 8B ternary MoE, 725B export. 3p×100 (MoE layers 3,16,22,29): one seed each; every final patch fails the +0.003 code-damage gate (FSA +0.004 / +0.011 / +0.007 at 1M / 4M / 16M, LoRA +0.008). Recipes are tuned per method (rehearsal, penalty, steps differ), so this compares tuned recipes rather than isolating the method.